

Introduction to Quantified Cognition

Psychology 5332 — Spring 2020

Per B. Sederberg, PhD

Version 2020-01-15

Quick Reference

Credit: 3 units ; Class # 19450

Time: Thursday, 15:30 – 18:00

Place: Ruffner Hall, Rm. 173

Text: Assigned readings

Course Web Page: Campuswire / GitHub

Course assistants: None

Instructor: Dr. Per Sederberg

Office: Gilmer Hall, Rm. B011 (for now)

E-mail: pbs5u@virginia.edu

Lab Website: Computational Memory Lab

Office hours: Tuesday and Thursday, 2–2:30pm

Final: Project-based

Overview and Course Objectives

Much of science, and especially psychology and neuroscience, involves testing and updating verbal theories, which are often imprecise and under-specified. Only by quantitative modeling of our experimental results can we hope to make significant progress in understanding the mechanisms that underlie cognition. Furthermore, only by quantifying and defining knowledge via mathematical principles can we achieve the effective interdisciplinary communication necessary for combining approaches to make useful progress towards understanding a system as complex as the brain. This course aims to equip students with the skills to a) think more deeply about the mechanisms underlying observed neural and behavioral phenomena, improving scientific thinking, and b) develop computational models that enable more precise, quantifiably testable hypotheses, improving the scientific process.

Topics covered in the course will include: computer programming, probability theory, Bayesian statistics, computational modeling, model comparison, neuroscience methods, neural networks, and open science. Taught at a high level, yet with practical hands-on examples for every topic, this course will provide the foundation necessary to understand, develop, and compare mechanistic models of cognitive processes.

Computing Requirements

This is a computational class and all work will be performed on a computer and almost entirely with the Python programming language within Jupyter notebooks. You will need to bring a laptop running Windows, OSX, or Linux to every class.

There are some online tools that provide free hosting and execution for Jupyter notebooks (e.g., Google Collaboratory). These hosting services limit the computational resources, so you may want to run Jupyter notebooks directly on your computer. This will also allow you to incorporate these approaches into your own research more easily. Thus, my recommendation is that you install and use the Anaconda Python distribution for your OS.

We will spend time on the first day of class to ensure everyone has a functioning computer that will be able to run everything necessary for the course either via Google Collaboratory or locally.

Assistance

I am eager for you to get as much as possible from this course, so please feel free to come to me with any questions you have. That said, science is a team effort and in order to reduce duplication of questions and discussions, we will be using Campuswire for all class communication and discussions. Please do not email me unless there is an issue with Campuswire. If you'd prefer to have a one-on-one discussion it is possible to send direct messages in Campuswire. We will spend some time on the first day ensuring everyone is set up to use Campuswire. I will also have weekly office hours to which you are always welcome to come and have in-person discussions.

Schedule

The following is the general order of the topics covered in the course. Please note that sometimes we may cover multiple topics in a single lecture, or spend more than one lecture on a single topic, and this list is subject to modification at any time. That said, all major changes will also include an update to the syllabus, so it will remain a point of reference.

0. Overview
1. Programming and Principles of Open Science
2. Probability
3. Data as a Random Variable
4. Quantifying Uncertainty
5. Bayesian Regression
6. Bayesian Data Analysis
7. Cognitive Process Models
8. Interactive Model Exploration
9. Bayesian Fits of Process Models
10. Inference with Process Models
11. Bring Your Own Data (BYOD) Project

Readings

There is no textbook for this course, however, there will often be relevant readings that accompany each lesson. PDF versions of these will be shared along with the Jupyter notebooks and lectures for each class.

Evaluation

This is a graduate-level course, which means that much of the burden of staying motivated to learn is transferred to the student. As such, there will not be any in-class exams. Your final grade will depend on class participation (60%), which will include taking part in interactive problem solving, homework, and presenting on or leading a discussion of a relevant paper. There will also be a final project involving a mechanistic data analysis and associated write-up, worth the remaining 40%.